

Goals

- implement the model $f_{w,b}$ for linear regression with one variable
- implement and explore the cost function for linear regression with one variable.

1. Model Representation

Notation

Here is a summary of some of the notation you will encounter.

General		Python (if applicable)
Notation	Description	
a	scalar, non bold	
\mathbf{a}	vector, bold	
Regression		
\mathbf{x}	Training Example feature values	<code>x_train</code>
\mathbf{y}	Training Example targets	<code>y_train</code>
$x^{(i)}, y^{(i)}$	i_{th} Training Example	<code>x_i, y_i</code>
m	Number of training examples	<code>m</code>
w	parameter: weight	<code>w</code>
b	parameter: bias	<code>b</code>
$f_{w,b}(x^{(i)})$	The result of the model evaluation at $x^{(i)}$ parameterized by w, b : $f_{w,b}(x^{(i)}) = wx^{(i)} + b$	<code>f_wb</code>

Tools

In this assignment you will make use of:

- NumPy, a popular library for scientific computing
- Matplotlib, a popular library for plotting data

```
In [9]: import numpy as np
import matplotlib.pyplot as plt
```

Problem Statement

As in the lecture, you will use the motivating example of diabetes progression prediction. This assignment will use a simple data set with only two data points shown below. These two points will constitute our *data or training set*.

BMI	Diabetes progression
32.1	151
21.6	75

You would like to fit a linear regression model through these two points, so you can then predict diabetes progression for other patients - say, a patient with BMI = 30.5.

Please run the following code cell to create your `x_train` and `y_train` variables. The data is stored in one-dimensional NumPy arrays.

```
In [10]: # x_train is the input variable (BMI)
# y_train is the target (diabetes progression level)
x_train = np.array([32.1, 21.6])
y_train = np.array([151, 75])
print(f"x_train = {x_train}")
print(f"y_train = {y_train}")

x_train = [32.1 21.6]
y_train = [151 75]
```

Note: The course will frequently utilize the python 'f-string' output formatting described [here](https://docs.python.org/3/tutorial/inputoutput.html) (<https://docs.python.org/3/tutorial/inputoutput.html>) when printing. The content between the curly braces is evaluated when producing the output.

Number of training examples m

You will use m to denote the number of training examples. Numpy arrays have a `.shape` parameter.

`x_train.shape` returns a python tuple with an entry for each dimension. `x_train.shape[0]` is the length of the array and number of examples as shown below.

```
In [11]: # m is the number of training examples
print(f"x_train.shape: {x_train.shape}")
m = x_train.shape[0]
print(f"Number of training examples is: {m}")

x_train.shape: (2,)
Number of training examples is: 2
```

One can also use the Python `len()` function as shown below.

```
In [12]: # m is the number of training examples
m = len(x_train)
print(f"Number of training examples is: {m}")

Number of training examples is: 2
```

Training example x_i , y_i

You will use $(x^{(i)}, y^{(i)})$ to denote the i^{th} training example. Since Python is zero indexed, $(x^{(0)}, y^{(0)})$ is (32.1, 151) and $(x^{(1)}, y^{(1)})$ is (21.6, 75).

To access a value in a Numpy array, one indexes the array with the desired offset. For example the syntax to access location zero of `x_train` is `x_train[0]`. Finish the next code block below to get the i^{th} training example.

```
In [13]: i = 0 # Change this to 1 to see (x^1, y^1)

x_i = x_train[i]           # ith feature value
y_i = y_train[i]           # ith target value
print(f"(x^{i}), y^{i}) = ({x_i}, {y_i})")

(x^(0), y^(0)) = (32.1, 151)
```

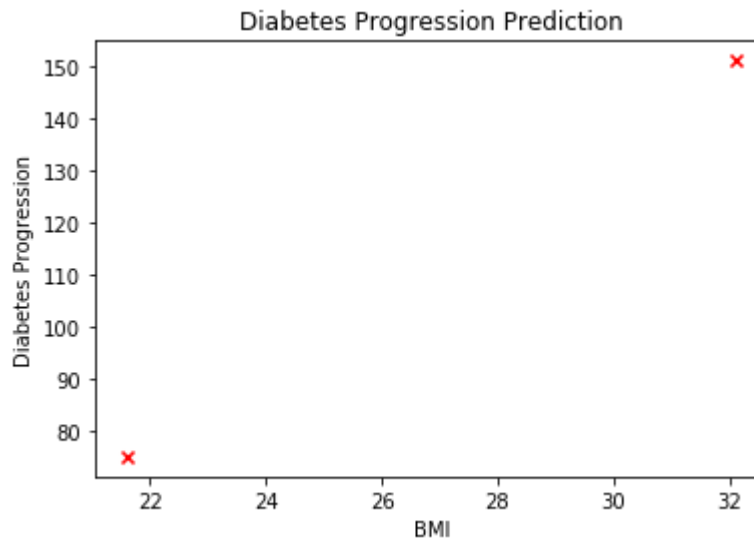
Plotting the data

You can plot these two points using the `scatter()` function in the `matplotlib` library, as shown in the cell below.

- The function arguments `marker` and `c` show the points as red crosses (the default is blue dots).

You can use other functions in the `matplotlib` library to set the title and labels to display

```
In [14]: # Plot the data points
plt.scatter(x_train, y_train, marker='x', c='r')
# Set the title
plt.title("Diabetes Progression Prediction")
# Set the y-axis label
plt.ylabel('Diabetes Progression')
# Set the x-axis label
plt.xlabel('BMI')
plt.show()
```



Model function

As described in lecture, the model function for linear regression (which is a function that maps from x to y) is represented as

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b \quad (1)$$

The formula above is how you can represent straight lines - different values of w and b give you different straight lines on the plot.

Let's try to get a better intuition for this through the code blocks below. Let's start with $w = 1$ and $b = 1$.

Note: You can come back to this cell to adjust the model's w and b parameters

```
In [82]: w = 7.25
b = -82
print(f"w: {w}")
print(f"b: {b}")
```

```
w: 7.25
b: -82
```

Now, let's compute the value of $f_{w,b}(x^{(i)})$ for your two data points. You can explicitly write this out for each data point as -

for $x^{(0)}$, $f_{wb} = w * x[0] + b$

for $x^{(1)}$, $f_{wb} = w * x[1] + b$

For a large number of data points, this can get unwieldy and repetitive. So instead, you can calculate the function output in a `for` loop in the `compute_model_output` function below.

Note: The argument description `(ndarray (m,))` describes a Numpy n-dimensional array of shape `(m,)`. `(scalar)` describes an argument without dimensions, just a magnitude.

Note: `np.zeros(n)` will return a one-dimensional numpy array with n entries

```
In [83]: def compute_model_output(x, w, b):
          """
          Computes the prediction of a linear model
          Args:
              x (ndarray (m,)): Data, m examples
              w,b (scalar)      : model parameters
          Returns
              y (ndarray (m,)): target values
          """
          m = x.shape[0]
          f_wb = np.zeros(m)
          # write a loop to compute f_wb
          for i in range(len(f_wb)):
              f_wb[i] = w * x[i] + b

          return f_wb
```

Now let's call the `compute_model_output` function and plot the output..


```
In [84]: tmp_f_wb = compute_model_output(x_train, w, b)
# call the compute_model_output function

# Plot our model prediction
plt.plot(x_train, tmp_f_wb, c='b', label='Our Prediction')

# Plot the data points
plt.scatter(x_train, y_train, marker='x', c='r')

# Set the title
plt.title("Diabetes Progression Prediction")
# Set the y-axis label
plt.ylabel('Diabetes Progression')
# Set the x-axis label
plt.xlabel('BMI')

plt.legend()
plt.show()
```



As you can see, setting $w = 1$ and $b = 1$ does *not* result in a line that fits our data.

Prediction

Try experimenting with different values of w and b . What should the values be for a line that fits our data? Note that you can actually compute the theoretical values of w and b by hand given the two training examples. Put your best w and b in the prediction cell below.

Now that we have a model, we can use it to make our original prediction. Let's predict the diabetes progression of a patient with BMI=30.5. Note: your prediction value should be around 140.

```
In [89]: w = 7.25          # Your best w
         b = -82           # Your best b
         x_i = 30.5
         diabetes_progression = w * x_i + b          # prediction

         print(f"${diabetes_progression:.1f}")

$139.1
```

2. Cost Function

Here, cost is a measure of how well our model is predicting the diabetes progression of a patient.

The equation for cost with one variable is:

$$J(w, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)})^2 \quad (2)$$

where

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b \quad (3)$$

- $f_{w,b}(x^{(i)})$ is our prediction for example i using parameters w, b .
- $(f_{w,b}(x^{(i)}) - y^{(i)})^2$ is the squared difference between the target value and the prediction.
- These differences are summed over all the m examples and divided by $2m$ to produce the cost, $J(w, b)$.

Note, in lecture summation ranges are typically from 1 to m , while code will be from 0 to $m-1$.

The code below calculates cost by looping over each example. In each loop:

- f_{wb} , a prediction is calculated
- the difference between the target and the prediction is calculated and squared.
- this is added to the total cost.

```
In [94]: def compute_cost(x, y, w, b):
        """
        Computes the cost function for linear regression.

        Args:
            x (ndarray (m,)): Data, m examples
            y (ndarray (m,)): target values
            w,b (scalar)      : model parameters

        Returns
            total_cost (float): The cost of using w,b as the parameters
            for linear regression
            to fit the data points in x and y
        """
        # number of training examples
        m = x.shape[0]

        cost_sum = 0

        # write a loop to compute the summation of the squared difference
        for all training examples

        cost_sum = sum([(w * x[i] + b) - y[i])**2 for i in range(m)])

        total_cost = (1 / (2 * m)) * cost_sum

        return total_cost
```

Test your compute_cost function using `x_train`, `y_train`, and your best w and b . Your total cost should be around 0.

```
In [95]: total_cost = compute_cost(x_train, y_train, w, b)
        # call compute_cost function
        print(f"Total cost is: {total_cost:.1f}")

        Total cost is: 0.1
```

Redefine your `x_train` and `y_train` using a larger training set below. Test your compute_cost function again.

BMI	Diabetes progression
32.1	151
21.6	75
30.5	141
22.6	97


```
In [97]: x_train = np.array([32.1, 21.6, 30.5, 22.6])
y_train = np.array([151, 75, 141, 97])
total_cost = compute_cost(x_train, y_train, w, b)
# call compute_cost function
print(f"Total cost is: {total_cost:.1f}")
```

Total cost is: 29.2

In []: